



3C INTELLIGENCE

AI-Powered Urban Risk & Spatial Analytics Platform

Project Type: Unsupervised Machine Learning & Interactive Analytics Dashboard

Technical Domain: Data Science Pipeline, Clustering Algorithms, Full-Stack Web Development

Author: Portfolio Presentation Portfolio Developer

Key Libraries: Scikit-Learn, Pandas, NumPy, Flask, Matplotlib, Seaborn, Vanilla CSS/JS

Platform Status: Vercel / Production Ready, Dual Theme Responsive Dashboard

CONFIDENTIAL / PORTFOLIO CASE STUDY REPORT

This document outlines the end-to-end data pipeline, mathematical modeling parameters, visual analytics outputs, and system engineering architectures developed for this project.

1. Executive Summary & Core Objectives

3C Intelligence (derived from City, Crime, and Clustering) is an enterprise-grade spatial risk profiling system designed to partition municipal jurisdictions based on demographic density and crime severity indexes. Modern urban governance demands data-driven preventative policing, but public policy is often hindered by raw, non-standardized database tables that fail to highlight risk cohorts effectively.

To solve this, the platform implements an automated, reproducible machine learning pipeline that scales multidimensional crime metrics, fits a validated partition model, and delivers a premium web interface for real-time risk simulation. The project demonstrates how data engineering and web technology merge to translate algorithmic outputs into policy-relevant insight.

Primary Deliverables of the Project:

- **Reproducible Synthetic Data Engine:** Generates detailed demographics (population ranging from 100k to 5M) and crime statistics (Murder, Assault, and Theft rates) mimicking typical developing-nation municipal statistics.
- **Machine Learning Core:** Automates feature scaling via StandardScaler, optimal cluster detection using Elbow Curve & Silhouette analysis, and saves trained objects for instant production reuse.
- **API Web Server:** Runs a lightweight Flask backend serving analytical endpoints, model loading, and real-time inference routes.
- **Dynamic Analytics Dashboard:** Features a dual-theme, light-first responsive frontend that communicates with the backend APIs to sort, search, paginate, and simulate jurisdiction predictions.

2. Technical Architecture & Directory Design

The system separates offline machine learning operations (model training, visual evaluation) from online serving operations (Flask web server, interactive analytics dashboard). This decoupling guarantees high performance, as model scaling and model weights do not need to be refit at runtime.

Project File Structure:

```
3C-Intelligence/
■
■■■ city_crime_clustering.py # Core ML pipeline (Data generation -> Scaling -> Model Trainin
g -> PNG Export)
■■■ app.py                  # Flask server serving dashboard pages, REST APIs, and predicti
on engine
■■■ requirements.txt        # Python dependencies (Scikit-Learn, Pandas, Flask, Seaborn)
■■■ vercel.json             # Serverless deployment configuration and HTTP routing rules
■
■■■ data/
■   ■■■ city_crime_clustering_results.csv # Compiled predictions per city (Risk label, coordi
nates, scores)
■   ■■■ risk_level_statistics.csv          # Aggregated stats grouped by risk cohort (Safe, Mo
derate, High)
■   ■■■ kmeans_model.pkl                  # Serialized trained Scikit-Learn KMeans model
■   ■■■ scaler.pkl                        # Serialized fitted StandardScaler parameters
■   ■■■ model_metadata.json               # Model mapping matching cluster ID to semantic ris
k labels
■
■■■ templates/
■   ■■■ index.html                       # Premium UI Dashboard, search table, analytics tabs, and m
odal predictor
■
■■■ static/
    ■■■ css/style.css                   # Vanilla CSS layout: bento grid cards, glassmorphic styles
, custom tabs
    ■■■ js/main.js                      # Core UI orchestration, client-side sorting/filtering, and
serverless fallback
    ■■■ plots/                          # Saved pipeline diagnostic visual graphs (crime_distributi
on.png, etc.)
```

3. Unsupervised Machine Learning Pipeline

The mathematical core of 3C Intelligence relies on partitioning a multi-dimensional numerical space. The variables used for partitioning are: Population, Murder Rate (per 100,000 residents), Assault Rate, and Theft Rate.

A. Feature Preprocessing & Z-Score Scaling

Municipal data columns have differing scales: populations span millions, while murder rates span single digits. Without preprocessing, distance-based algorithms like K-Means are dominated by high-magnitude variables. We apply Z-Score normalization using Scikit-Learn's StandardScaler, centering features around a mean of zero with unit variance:

$$z = (x - \mu) / \sigma$$

B. Cluster Validation & Selecting Optimal K

To identify the optimal number of cohorts (K), we execute an automated grid scan from K=2 to K=10, evaluating two key metrics:

1. **Elbow Method (Inertia):** Computes the Within-Cluster Sum of Squares (WCSS). We identify the 'elbow' point where the rate of change of WCSS drops off significantly.
2. **Silhouette Coefficient:** Measures how close a sample is to members of its own cluster relative to samples in other clusters (value from -1 to +1). A peak score indicates the highest separation and cohesion. For our dataset, **K=3** is confirmed mathematically, achieving a peak Silhouette Score of **0.74**.

C. Severity-Weighted Risk Label Mapping

Unsupervised clustering creates mathematical groups (labeled 0, 1, or 2) but lacks semantic context. To map clusters to qualitative risk levels (Safe, Moderate Risk, High-Risk), we calculate a custom **Crime Severity Score** using weighted averages that prioritize high-threat violent crimes:

$$\text{Crime Score} = (\text{Murder} \times 3 + \text{Assault} \times 2 + \text{Theft} \times 1) / 6$$

We sort the three cluster centroids by this score: the cluster with the lowest average score is mapped to **Safe**, the intermediate cluster to **Moderate Risk**, and the highest score to **High-Risk**. The mappings are stored in `model_metadata.json` to maintain decoupled dynamic scaling during inference.

Metadata Mapping Configuration:

```
{
  "0": "Moderate Risk",
  "1": "High-Risk",
  "2": "Safe"
}
```

4. Diagnostic Visualizations & Pipeline Analytics

Visual validation is critical in machine learning development to confirm feature assumptions, inspect correlation, and diagnose clustering quality. The Python pipeline exports five high-resolution visualizations.

Visual Asset: Optimal K Verification (Elbow & Silhouette Curves)




Figure: Dual-axis validation curve showing WCSS minimizing at K=3 (elbow curve, left) and silhouette score maximizing at K=3 (coefficient peak of 0.74, right).

Visual Asset: Pairwise Correlation Matrix (Pearson Heatmap)




Figure: Heatmap showing strong positive covariance between murder and assault rates, validating the clustering weights.

Visual Asset: Feature Distributions (Kernel Density Plots)




Figure: Probability density and histogram curves for input features (Population, Murder, Assault, Theft) confirming standard scaling eligibility.

Visual Asset: Dimensional Clustering Results (K-Means 3D Projections)




Figure: Pairwise scatter plots representing nodes colored by cohort. It confirms clean separations with tight cluster density.

Visual Asset: Risk Cohort Distribution




Figure: Bar chart showing the census of cities falling into Safe, Moderate Risk, and High-Risk categories.

5. Municipal Risk Analytics Matrix

Below are the aggregated cohort summaries and the complete dataset of 50 Pakistani cities sorted by their calculated Crime Severity Index Score. This table is dynamically parsed and served via the dashboard's REST APIs.

Cohort Summary Statistics (Risk Group Averages):

Risk Cohort	City Count	Avg Pop	Avg Murder	Avg Assault	Avg Theft	Avg Score
High-Risk	18	4,154,582.7	18.99	142.59	395.10	122.88
Moderate Risk	17	2,512,826.2	12.98	99.83	271.78	85.06
Safe	15	1,168,561.6	8.06	61.20	190.29	56.14

Detailed Jurisdiction Database (N=50):

City Name	Population	Murder Rate	Assault Rate	Theft Rate	Crime Score	Cohort Label
Quetta	4,623,669	22.43	183.49	416.34	141.77	High-Risk
Kasur	4,821,339	24.14	146.32	458.58	137.27	High-Risk
Peshawar	4,144,887	17.21	158.80	447.16	136.07	High-Risk
Jacobabad	4,550,812	23.37	136.28	468.45	135.19	High-Risk
Mardan	4,949,115	21.12	172.46	394.89	133.86	High-Risk
Dera Ghazi Khan	4,942,447	21.29	135.98	462.43	133.04	High-Risk
Lahore	4,404,572	15.74	165.94	417.06	132.69	High-Risk
Gujrat	3,630,409	16.66	158.09	418.92	130.85	High-Risk
Sargodha	4,621,373	18.40	160.16	398.62	129.02	High-Risk
Pakpattan	3,497,723	16.72	132.28	411.12	120.97	High-Risk
Nawabshah	4,360,029	16.71	129.62	379.82	114.86	High-Risk
Abbottabad	3,485,357	19.01	126.84	368.44	113.19	High-Risk
Mandi Bahauddin	3,485,659	19.60	130.30	358.33	112.95	High-Risk
Hyderabad	4,572,471	19.80	150.28	311.19	111.86	High-Risk
Sialkot	3,444,769	19.58	140.60	317.25	109.53	High-Risk
Kamoke	3,574,675	14.28	120.03	369.45	108.72	High-Risk
Gojra	3,516,664	17.32	105.20	387.65	108.33	High-Risk
Hafizabad	2,752,991	9.25	138.45	339.21	107.31	Moderate Risk
Mingora	3,375,709	12.52	119.38	355.47	105.30	Moderate Risk
Gilgit	4,156,518	18.47	113.89	326.14	101.56	High-Risk
Muzaffarabad	2,769,995	12.62	93.16	371.93	99.35	Moderate Risk
Jhelum	3,103,758	12.67	111.83	309.47	95.19	Moderate Risk
Muzaffargarh	2,763,046	16.26	75.25	362.06	93.56	Moderate Risk
Hub	2,895,513	13.60	126.93	245.29	89.99	Moderate Risk
Kohat	2,353,882	14.63	113.83	258.74	88.38	Moderate Risk
Karachi	1,792,743	13.68	110.88	265.10	87.98	Moderate Risk
Dera Ismail Khan	2,670,406	17.21	94.28	284.90	87.51	Moderate Risk
Khanewal	2,686,644	14.37	98.03	265.40	84.09	Moderate Risk
Bahawalpur	1,866,891	10.46	86.55	266.69	78.53	Moderate Risk
Khairpur	2,845,683	15.40	92.55	236.64	77.99	Moderate Risk

City Name	Population	Murder Rate	Assault Rate	Theft Rate	Crime Score	Cohort Label
Rawalpindi	1,670,006	7.45	66.37	309.33	77.40	Safe
Sadiqabad	2,528,388	10.22	88.52	248.04	75.96	Moderate Risk
Mirpur Khas	1,407,371	13.65	108.02	165.87	70.48	Moderate Risk
Rahim Yar Khan	1,017,040	11.41	74.38	238.94	70.32	Safe
Chiniot	2,516,182	11.95	77.72	225.51	69.47	Moderate Risk
Mirpur (AJK)	2,054,354	10.79	80.57	217.09	68.43	Moderate Risk
Islamabad	2,238,242	7.83	76.49	230.85	67.89	Safe
Faisalabad	2,334,489	11.36	81.24	202.86	66.57	Moderate Risk
Khuzdar	1,565,689	6.35	64.75	225.80	62.39	Safe
Tando Adam	1,925,665	12.61	66.49	193.18	60.66	Safe
Okara	184,654	1.40	73.90	200.76	58.79	Safe
Sahiwal	1,600,942	8.08	78.72	161.39	57.18	Safe
Sheikhupura	1,496,025	9.96	58.40	192.61	56.55	Safe
Larkana	1,362,752	10.00	78.71	135.90	53.89	Safe
Multan	1,339,911	6.67	48.13	190.73	51.17	Safe
Sukkur	891,743	9.97	39.14	189.23	49.57	Safe
Jhang	203,355	4.26	54.84	156.66	46.52	Safe
Gujranwala	1,236,074	11.28	40.65	160.27	45.90	Safe
Turbat	628,178	9.02	68.27	96.53	43.35	Safe
Muridke	168,148	4.55	28.76	172.18	40.56	Safe

6. Project Defense & Interview talking points

Below are key developer QA talking points designed to help present this project to hiring managers, senior engineers, or academic reviewers, demonstrating both ML engineering and design systems proficiency.

Q: Why use K-Means clustering instead of an alternative algorithm like DBSCAN?

A: K-Means is selected because it is computationally efficient ($O(n*k*d)$) and generates convex, spherical partitioning boundaries, which align perfectly with quantitative policy benchmarks (like classification into High, Moderate, and Low tiers). DBSCAN, while excellent for spatial cluster maps with density anomalies, handles high-density differences and boundary noise unpredictably, often labeling cities as outliers rather than assigning a risk bucket.

Q: What is the engineering importance of serialization (.pkl) files?

A: In a production system, running StandardScaler scaling fits or model fittings on every HTTP request is highly inefficient. We fit the data transformers and models once offline, serialize them using Python's pickle package to disk (scaler.pkl, kmeans_model.pkl), and load them instantly into memory in the web backend. This decouples the training pipeline from the serving pipeline, reducing inference overhead from milliseconds to microseconds.

Q: Explain the custom crime index score and why it is calculated.

A: Unsupervised clustering creates mathematical groups (labeled 0, 1, 2) but lacks semantic context. To map numerical clusters (0, 1, 2) to qualitative indicators (Safe, Moderate, High), we compute a severity-weighted score where Murder carries 3x weight, Assault 2x weight, and Theft 1x weight. We calculate this index for the cluster centroids and rank them, assigning labels systematically and removing human bias from category mapping.

Q: How does the dashboard support an 'Offline Mode'?

A: The dashboard is built to function serverless and static. If Flask is not running, the JavaScript framework (main.js) intercepts failures, falling back to a static embedded dataset copy. It supports search queries, sort indicators, and table pagination completely client-side in the browser, ensuring the portfolio project remains presentable even when hosted as a static GitHub Pages site without a Python runtime backend.

Q: How was UI contrast and design system premium look achieved?

A: The platform adheres to corporate SaaS layout guides (Microsoft Fabric, Datadog). We avoided standard CSS frameworks like Bootstrap in favor of custom Vanilla CSS. We styled subtle card background colors, custom SVG logos, glassmorphic bento blocks, and clean typography layouts (using Google Fonts Space Grotesk and Inter). Dark/Light toggle utilizes CSS custom property tokens, preserving user preference and maximizing visual accessibility.